

**Desmond Oketch- SCT212-0083/2021**

## **COMPUTER TECHNOLOGY**

### **Lab 5**

#### **Computer Architecture – tutorial 5**

#### **Problem Summary**

Analyse the data cache behaviour for the given code:

```
for (i = 0; i < 4096; i++)
```

```
    X[i] = X[i] * Y[i] + C;
```

Arrays:

- X[4096] and Y[4096] → 4096 elements of 4 bytes = 16 KB each.

Cache:

- Size: 16 KB
- Line size: 64 Bytes
- Associativity: Direct-mapped
- Addressing: Physical

Compute total cache lines:

Cache lines =  $\frac{16 \text{ KB}}{64 \text{ B}} = 256 \text{ lines}$   
 $\text{Cache lines} = \frac{16384}{64} = 256 \text{ lines}$   
Cache lines =  $\frac{64 \text{ KB}}{16 \text{ B}} = 256 \text{ lines}$

### **Part a: Cache Miss Analysis**

#### **Memory Access Pattern**

Each iteration:

1. Load X[i] → 4 bytes
2. Load Y[i] → 4 bytes
3. Store X[i]

Total: 3 memory operations per iteration.

Each cache line holds:

$$64 / 4 = 16 \text{ elements} \quad \frac{64}{4} = 16 \text{ elements}$$

So, accessing  $X[i]$  will bring  $X[i]$  to  $X[i+15]$  into cache.

## Compulsory Misses (cold misses)

Each unique cache line accessed for the first time causes a compulsory miss.

X accesses:

- $4096 \text{ elements} / 16 \text{ per line} = 256 \text{ unique cache lines}$
- So, 256 compulsory misses.

Y accesses:

- $4096 \text{ elements} / 16 \text{ per line} = 256 \text{ unique lines}$

## Conflict Misses

The cache is direct-mapped and X & Y are stored consecutively, they may map to the same cache lines.

- Let's assume X starts at address 0.
- Since Y is allocated right after X, it will map to the same cache lines as X.

For every Y access, it evicts X from the cache.

Thus:

- When you load  $X[i]$ , it's a miss (first time: compulsory, later: conflict).
- Then you load  $Y[i] \rightarrow$  also a miss, replaces  $X[i]$ .
- Then you store to  $X[i] \rightarrow$  miss again, because  $X[i]$  was evicted by  $Y[i]$ .

So each iteration causes 3 cache misses:

- 1 for X[i] (load)
- 1 for Y[i]
- 1 for X[i] (store)

## Total Misses

$4096 \times 3 = 12288$  data cache misses

## Breakdown

- Compulsory Misses:  $256 (X) + 256 (Y) = 512$
- Conflict Misses:  $12288 - 512 = 11776$
- Capacity Misses: 0, because total working set = 32KB, but we're not accessing more than 2 elements at once.

## Cache Miss Rate

Total memory accesses =  $4096 \times 3 = 12288$

Miss rate =  $\frac{12288}{12288} = 100\%$

## Part b: Software Optimization

### Reduce conflict misses.

Solution: Loop Interchange or Blocking, but simpler here:

Interleave accesses to avoid eviction:

Instead of accessing X then Y:

```
for (i = 0; i < 4096; i++) {
    tempX = X[i]; // load X
    tempY = Y[i]; // load Y
    tempX = tempX * tempY + C;
    X[i] = tempX; // store X
}
```

## Better Software Fix: Loop Blocking

Use blocking with a block size that fits within the cache without causing conflicts.

```
int B = 64; // block size (or 128)
for (ii = 0; ii < 4096; ii += B) {
    for (i = ii; i < ii + B; i++) {
        X[i] = X[i] * Y[i] + C;
    }
}
```

## Best fix: Padding arrays

Add padding between X and Y so they don't map to same cache lines.

For example:

```
float X[4096];
float pad[256]; // adds 1KB of space
float Y[4096];
```

This makes sure that X and Y don't conflict in cache.

## Cache Misses Now:

- X accesses: 256 loads + 256 stores = 512 misses (compulsory only)
- Y accesses: 256 loads = 256 misses (compulsory)

Total = 768 cache misses

## Breakdown:

- Compulsory Misses: 768
- Conflict Misses: 0 (due to padding)
- Capacity Misses: 0

## Miss Rate

Total memory accesses = 12288

Miss rate =  $\frac{768}{12288} = 6.25\%$

## Part c: Hardware Optimization

### Modify the hardware

Solution: Change Cache Organization

Change from direct-mapped  $\rightarrow$  2-way set associative or fully-associative.

2-way set associative cache:

- $16\text{KB} / 64\text{B} = 256$  lines  $\rightarrow$  128 sets of 2 lines
- X and Y can now both map to the same set but occupy different ways  $\rightarrow$  no conflict.

### Cache Misses:

- First access of each block: compulsory
- No conflict misses (because associativity allows both to reside)

Same result as software padding.

### Total:

- X: 256 load + 256 store = 512
- Y: 256 load = 256

$\rightarrow$  Total = 768

### Breakdown:

- Compulsory: 768
- Conflict: 0
- Capacity: 0

**Miss Rate:**

$$\text{Miss rate} = \frac{768}{12288} = 6.25\% \quad \frac{768}{12288} = \boxed{6.25\%} \quad \frac{768}{12288} = 6.25\%$$